

Time Copilot Prototype: Two-Page Business and Modeling Report

Executive Summary

This prototype tests whether a Time Copilot workflow can reliably select high-performing time-series models on open market data, with outputs that are actionable for trading and operations decisions. The result is yes: auto-selected champions consistently beat naive baselines and stay competitive against strong public baselines.

Key outcomes: - Forecast accuracy improved materially versus naive in both markets (PJM: -34.6% sMAPE, NP: -15.3% sMAPE). - Rally-risk prediction quality improved sharply versus naive (PJM: 2.31x PR-AUC, NP: 2.43x PR-AUC). - Winners differed by market and task, validating the need for automatic per-task/per-market model selection instead of one global model.

1) Business Problem (Why This Matters)

Commodity-facing teams make repeated decisions under uncertainty: - **Level decisions:** what is the most likely near-term value (price, spread, throughput proxy)? - **Risk decisions:** what is the probability of an adverse or opportunity event (rally/spike/outage-like condition)?

Traditional setups usually optimize one model for one metric and then reuse it broadly. This misses two realities: - Different markets have different dynamics. - Different decision types require different objective functions.

This prototype addresses both by benchmarking model families across two tasks and automatically selecting the best model by task and market.

2) Data Lineage and Governance

Data lineage was intentionally kept explicit and reproducible.

| Step | Artifact | What Happens | Leakage / Repro Control |
|--------|-----------------------------------|--|--|
| Source | EPF open datasets on Zenodo | Raw market data loaded for PJM and NP | Source is public and version-addressable |
| Cache | datasets/PJM.csv, datasets/NP.csv | Data cached locally for deterministic reruns | Same cached input reused across runs |

| Step | Artifact | What Happens | Leakage / Repro Control |
|---------------|---|--|--|
| Normalize | <code>normalize_epf_frame</code> | Canonical schema: <code>timestamp</code> , <code>price</code> | Timestamp parsing and schema validation |
| Feature build | <code>build_features</code> | Lag/rolling/calendar features (<code>lag_1</code> , <code>lag_24</code> , etc.) | Rolling stats are shifted to avoid look-ahead |
| Labeling | <code>build_rally_labels</code> and <code>label_future_rally</code> | Rally event target and future-horizon probability target | Future target is explicit and horizon-controlled |
| Split | <code>_train_test_split</code> | Chronological split (80/20) | Time-ordered split prevents train/test overlap |
| Evaluation | <code>forecast_metrics</code> , Forecast: <code>classification_metrics</code> | MAE/RMSE/sMAPE contracts across Rally: PR-AUC/ROC-AUC/F1/Brier | Shared metric |
| Selection | <code>pick_forecast_champion</code> and <code>pick_rally_champion</code> | pick best model by task | Rule-based, deterministic tie-breakers |

Primary output artifacts: - `artifacts/dual_market/<market>/forecast_benchmark.csv`
- `artifacts/dual_market/<market>/rally_benchmark.csv` - `artifacts/dual_market/champion_summary.csv`

3) Benchmark Setup and Model Set

Two benchmark tasks: - **Forecast task**: next-step price prediction (optimize lower sMAPE). - **Rally task**: probability of rally within the next 24-hour horizon (optimize higher PR-AUC).

Model families: - **Reference**: `naive` - **Strong public baselines**: `lear` (regularized linear autoregressive), `logreg` (balanced logistic regression) - **Nonlinear learners**: `gbdt_reg`, `gbdt_cls` - **Neural**: `dnn_reg`, `dnn_cls`

Auto-pick rules: - Forecast champion = minimum sMAPE (tie-break: MAE) - Rally champion = maximum PR-AUC (tie-break: minimum Brier)

4) Performance Results (Including Strong Baselines)

Champion outcomes by market/task

| Market | Task | Champion | Primary Metric | Champion Value |
|--------|----------|-----------------------|---------------------------|----------------|
| PJM | Forecast | <code>gbdt_reg</code> | sMAPE (lower is better) | 12.01 |
| PJM | Rally | <code>gbdt_cls</code> | PR-AUC (higher is better) | 0.870 |
| NP | Forecast | <code>dnn_reg</code> | sMAPE (lower is better) | 5.54 |
| NP | Rally | <code>dnn_cls</code> | PR-AUC (higher is better) | 0.789 |

Lift versus naive and versus strongest non-champion baseline

| Market | Task | Lift vs Naive | Strong Baseline Comparator | Lift vs Strong Baseline |
|--------|----------|---------------------------------------|---------------------------------|-------------------------|
| PJM | Forecast | 34.6% lower sMAPE (18.36 -> 12.01) | <code>dnn_reg</code> (13.20) | 9.0% lower sMAPE |
| NP | Forecast | 15.3% lower sMAPE (6.54 -> 5.54) | <code>gbdt_reg</code> (5.62) | 1.4% lower sMAPE |
| PJM | Rally | +0.493 PR-AUC (2.31x) | <code>dnn_cls</code> (0.853) | PR-AUC (+2.0%) |
| NP | Rally | +0.464 PR-AUC (2.43x) | <code>logreg</code> (0.787) | PR-AUC (+0.26%) |

Interpretation: - Gains over naive are large and consistent. - Against strong baselines, gains are market/task dependent but still positive; in NP rally, improvement is marginal, which is exactly the kind of nuance an auto-pick system should surface.




Figure 1: Champion scorecard




Figure 2: Forecast sMAPE by model




Figure 3: Rally PR-AUC by model




Figure 4: Forecast efficiency frontier




Figure 5: Rally calibration curves

5) Chart Evidence (Executive + Analytical)

Champion scorecard

Forecast benchmark (sMAPE)

Rally benchmark (PR-AUC)

Forecast efficiency frontier (error vs RMSE)

Rally calibration (top model per market)

6) Why This Changes Time-Series Modeling Practice

This prototype is important less because one model wins, and more because it changes **how** modeling is operationalized:

- Moves from single-model forecasting to **portfolio benchmarking**.
- Separates **level prediction** from **event-risk prediction**, matching real decision needs.
- Treats model selection as a governed, repeatable output (`champions.json`), not ad hoc analyst choice.
- Makes model swaps cheap: adding/removing models does not change the evaluation contract.
- Produces decision-ready artifacts (tables + charts) that non-ML stakeholders can review quickly.

In short, this is a shift from “fit one time-series model” to “run a model selection system for time-series decisions.”

7) Limits and Next Validation Gate

Current limits:

- Public electricity data is a proxy, not your internal production domain.
- Single chronological holdout (80/20) is useful for prototype speed but not the final governance standard.
- Neural models run, but further hyperparameter tuning could improve stability and consistency.

Next gate before internal rollout:

1. Re-run the exact framework on internal targets with the same contracts and metrics.
2. Add rolling backtests (multiple time windows) to validate stability.
3. Compare against the desk’s current heuristic or incumbent model as the operational baseline.